


Van Hateren disk teacher: full 10,000 pixels, n=1,000



MC: arithmetic mean of 100 trials; error bars ± 2 SE; shading: trial 10-90%, not confidence intervals.

*Gaussian surrogate: 8,192 draws; R^2 acc uses per-draw inverse slopes (not transformed mean slopes).

Middle column only: triangles/dotted lines = empirical sklearn LOOCV; hatched bands = its 10-90% trials.

Open squares: direct formula at boundary κ_0 (not stationary); $\lambda^*=0$ omitted on log axis.